

Table 1: Ghost-skill structure and fold-stability summary. $\text{Skill}_{\text{RMSE}}$ and variance retention are pooled over matching fold-horizon windows. Fold-wise variance retention: median [min, max] across 5 folds, reported only for $h \geq 24$ h where fold heterogeneity is material. Source: `pub_table_4_ghost_skill_structure.csv` and `fold_stability_summary_sarima.csv`, experimental commit 95c9cbd.

Model	h (h)	$\text{Skill}_{\text{RMSE}}$	Var. ret. (pooled)	POD	CSI	Ghost-skill status
LightGBM	1	−0.0104	0.689	0.773	0.658	Negative skill
LightGBM	6	0.0729	0.551	0.585	0.450	Moderate degradation; events retained
LightGBM	24	−0.0343	0.599	0.472	0.328	Negative skill
LightGBM	48	−0.0793	0.866	0.323	0.214	Negative skill
SARIMA	1	0.0314	0.798	0.785	0.645	No ghost skill
SARIMA	6	0.0827	0.357	0.534	0.425	No ghost skill
SARIMA	24	0.0450	0.029	0.099	0.095	Strong candidate; fold heterogeneity
SARIMA	48	0.1325	0.004	0.000	0.000	Ghost skill: 3/5 folds

Note: At $h = 48$ h, SARIMA fold-wise variance retention: median 0.0007 [0.07%], range [0.0005, 0.0012]. Dynamic collapse in 5/5 folds; complete event failure in 5/5 folds; positive error skill in 3/5 folds; full ghost-skill pattern in 3/5 folds.